

# AI Salesman System Documentation

## Overview

The AI Salesman system is designed to act as a cold-calling agent capable of holding meaningful conversations, understanding customer queries, and closing sales. The focus is on creating a high-performance, low-latency solution that delivers natural-sounding voice interactions, simulating human-like conversation.

The web-based demonstration showcases the system's functionalities without telephony integration. The solution uses advanced AI models for STT (Speech-to-Text), TTS (Text-to-Speech), and conversation handling with a robust backend to manage data and interactions.

---

## Architecture

### Workflow:

1. **Customer Call Initiation:** The AI Salesman begins a simulated call or receives one via the web interface.
  2. **Speech-to-Text (STT):** Real-time transcription of customer speech using Whisper Large V3 Turbo.
  3. **LLM (Large Language Model):** LLaMA 3.3 70B processes the transcribed text, decodes customer queries, and retrieves relevant data.
  4. **Context Management:** Tracks and stores the conversation context for personalized interactions and smooth conversation flow.
  5. **Data Retrieval:** Queries the Pinecone database for customer information and related details.
  6. **Chatbot Response:** Generates responses to customer queries.
  7. **Text-to-Speech (TTS):** Converts the chatbot's response to speech using F5 TTS (running locally).
  8. **Call Summary:** At the end of the call, a summary is generated and presented to the operator.
- 

### Technologies Used:

- **LLM:** LLaMA 3.3 70B
  - **STT:** Whisper Large V3 Turbo
  - **TTS:** F5 TTS (local deployment)
  - **Backend:** FastAPI
  - **Database:** Pinecone
  - **Web Demo:** HTML, CSS, and JavaScript frontend
- 

## Setup Instructions

### Prerequisites

1. Python 3.10+
2. Packages in requirements.txt
3. Docker (optional for containerized deployment)
4. F5 TTS installed locally
5. API keys for required services:
  - PINECONE\_API
  - LLAMA\_API\_KEY
  - HUGGING\_FACE\_API
  - SAMBANOVA\_API\_KEY

### Installation Steps:

#### Backend:

1. Clone the repository.
2. Install dependencies:

```
pip install -r requirements.txt
```
3. Set up environment variables in a `.env` file:

```
PINECONE_API=your_pinecone_api_key
LLAMA_API_KEY=your_llama_api_key
HUGGING_FACE_API=your_hugging_face_api_key
SAMBANOVA_API_KEY=your_sambanova_api_key
```
4. Start the backend server:

```
uvicorn web.main:app --host 0.0.0.0 -port $PORT
```

#### Frontend:

1. Navigate to the `web` directory.
2. No additional dependencies required for HTML, CSS, and JavaScript.
3. Run the `uvicorn` command then go to the browser to run the web demo.

#### F5 TTS:

1. Ensure F5 TTS is set up locally and running.
2. Configure the backend to communicate with the local TTS instance.

---

## File-Specific Details

### `main.py`

- **Purpose:** Serves as the entry point for the backend system.
- **Key Functions:**
  - Defines API routes for handling transcription requests (`/ws/speech-to-text`), chatbot interaction (`/ws/generate-response`), and TTS conversion.
  - Implements the FastAPI framework to ensure asynchronous and scalable request handling.
- **Implementation Details:**
  - Uses `uvicorn` for running the server.
  - Routes link directly to services like `stt_service.py` and `tts_service.py`.

### `stt_service.py`

- **Purpose:** Processes real-time audio input to transcribe customer speech.
- **Key Functions:**
  - Utilizes Whisper Large V3 Turbo for accurate transcription.
  - Handles audio stream preprocessing and manages API calls to the Whisper model.
- **Implementation Details:**
  - Ensures low-latency transcription for real-time interactions.
  - Converts audio input into text format compatible with the chatbot.

### `chatbot_service.py`

- **Purpose:** Handles natural language processing and response generation.
- **Key Functions:**
  - Integrates LLaMA 3.3 70b to generate contextually accurate and conversationally relevant responses.
  - Maintains conversation history and context for a seamless user experience.

- **Implementation Details:**
  - Supports multi-turn conversations.
  - Manages errors like unclear inputs with fallback responses.

### database\_service.py

- **Purpose:** Interfaces with Pinecone to store and retrieve customer data and conversation logs.
- **Key Functions:**
  - Manages vectorized data storage for efficient querying.
  - Retrieves context-specific information to personalize responses.
- **Implementation Details:**
  - Utilizes Pinecone's API for high-speed data indexing and lookup.
  - Ensures secure and optimized database operations.

### tts\_service.py

- **Purpose:** Converts text responses from the chatbot into natural-sounding audio.
- **Key Functions:**
  - Leverages the locally hosted F5 TTS engine for speech synthesis.
  - Processes chatbot output and generates audio streams for playback.
- **Implementation Details:**
  - Configurable to adjust voice pitch, speed, and tone.
  - Communicates with the backend through defined API routes.

### frontend/

- **Purpose:** Provides the user interface for the web-based demonstration.
  - **Components:**
    - **HTML:** Structures the layout for the web interface.
    - **CSS:** Designs the visual style and user-friendly aesthetics.
    - **JavaScript:** Implements interactive features like initiating calls, displaying transcripts, and audio playback.
  - **Implementation Details:**
    - Ensures cross-browser compatibility.
    - Uses asynchronous requests to interact with the backend endpoints seamlessly.
- 

## Features

1. Real-time transcription of customer speech.
2. Intelligent query handling with a state-of-the-art LLM.

3. Context-aware conversations for personalized interaction.
  4. Natural-sounding TTS output.
  5. Automated call summary generation.
- 

## API References

### STT Endpoint

- **Route:** `/ws/speech-to-text`
- **Method:** POST
- **Input:** Audio stream
- **Output:** Transcribed text

### LLM Endpoint

- **Route:** `/ws/generate-response`
- **Method:** POST
- **Input:** Transcribed text, conversation context
- **Output:** LLM response text

### TTS Endpoint

- **Route:** `/ws/generate-response`
  - **Method:** POST
  - **Input:** Chatbot response text
  - **Output:** Audio stream
- 

## Demonstration

Visit the web-based demo at [AI Salesman Demo](#) to experience the system in action.

---

## Future Enhancements

1. Integration with telephony systems for real-world deployment.
2. Advanced sentiment analysis for better customer interaction.
3. Additional language support for broader applicability.

---

## Contributors

- Devesh Bhushan, Devdeep Mukherjee, Dishita Pokharna, Bhumika

---

For further questions, please contact the development team.